

Regression Examples

October 11, 2021

```
[ ]: # import all the required libraries and put matplotlib in inline mode to plot
      ↪ on the notebook
import pandas as pd
import numpy as np
import math
import matplotlib.pyplot as plt
%matplotlib inline

# linear regression
from sklearn import linear_model

# nearest neighbor
from sklearn import neighbors

# regression trees (simple and ensemble)
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import AdaBoostRegressor

from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
```

```
[ ]: # define the random seed if needed
random_seed = 1234

# define the figure size and the font size
fig_width = 12
fig_height = 9
fig_font_size = 16
```

```
[ ]: df = pd.read_csv('housing.csv')
df.describe(include='all')
```

```
[ ]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500

50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000

	AGE	DIS	RAD	TAX	PTRATIO	B \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500
50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000

	LSTAT	MEDV
count	506.000000	506.000000
mean	12.653063	22.532806
std	7.141062	9.197104
min	1.730000	5.000000
25%	6.950000	17.025000
50%	11.360000	21.200000
75%	16.955000	25.000000
max	37.970000	50.000000

```
[ ]: def plot_prediction(x,y,xp=None,yp=None,title=""):
    """Plots the original data (x,y) and a set of point (xp,yp) showing the
    ↪ model approximation"""

    plt.figure(figsize=(fig_width,fig_height))

    font = {'family' : 'sans', 'size' : fig_font_size}
    plt.rc('font', **font)

    plt.scatter(x, y, color='blue')

    if (not xp is None) and (not yp is None):
        plt.plot(xp, yp, color='red', linewidth=3)

    plt.xlabel("LSTAT")
    plt.ylabel("MEDV")

    plt.title(title)

    plt.xlim([0,40])
    plt.ylim(0,55)

    plt.show()
```

```
[ ]: # compute the data inputs
X = df['LSTAT'].values.reshape(-1,1)
y = df['MEDV'].values.ravel()

# input values used to plot the predicted model
X_test = np.linspace(np.min(X), np.max(X), 500)[:, np.newaxis]

[ ]: plt.figure(figsize=(fig_width,fig_height))

font = {'family' : 'sans', 'size' : fig_font_size}
plt.rc('font', **font)


plt.scatter(X, y, color='blue')

plt.xlabel("LSTAT")
plt.ylabel("MEDV")


plt.title("Boston Housing Data (LSTAT vs MEDV)")

plt.xlim([0,40])
plt.ylim(0,55)

plt.show()
```




```
[ ]: plot_prediction(X,y);
```



```
[ ]: # simple linear regression
regr = linear_model.LinearRegression()
regr.fit(X, y)

# model output for the input data
y_test = regr.predict(X_test)

# plot the result
plot_prediction(X, y, X_test, y_test, "y=%.3fx + %.3f" % (regr.coef_[0], regr.
↪ intercept_));
```

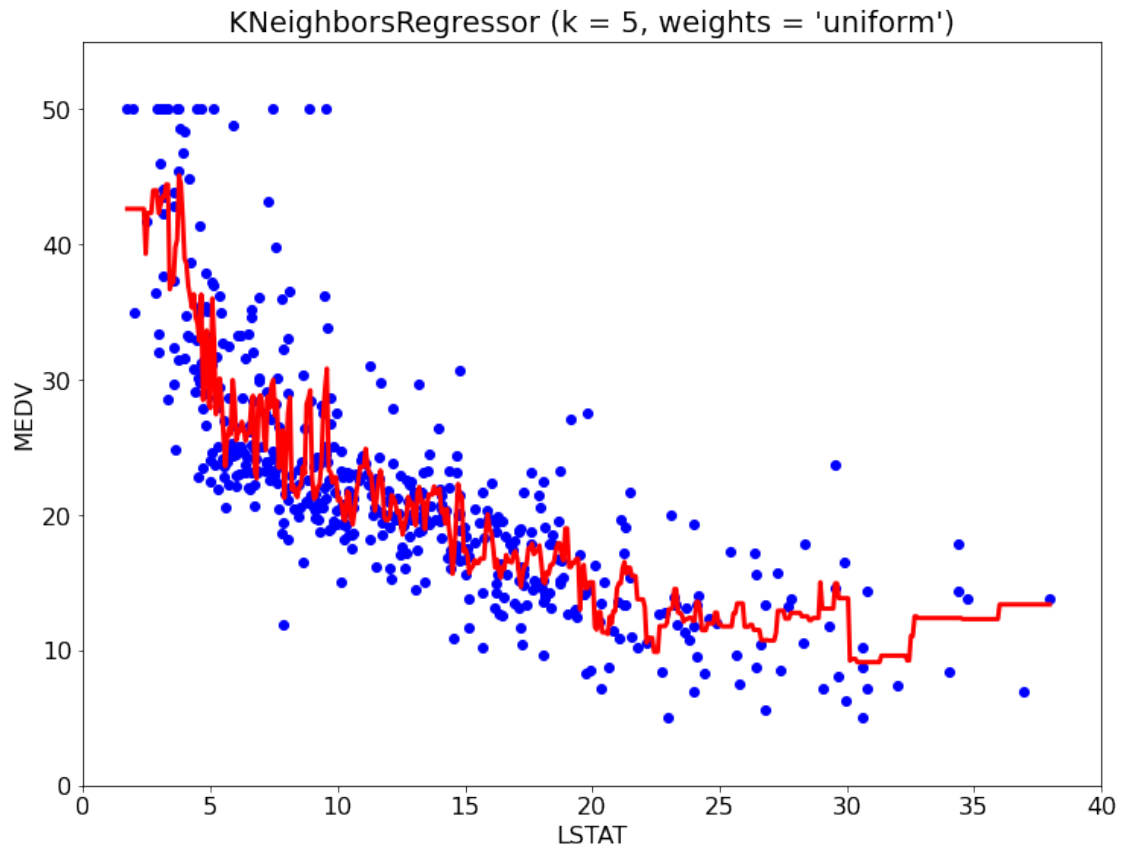


```
[ ]: knn = neighbors.KNeighborsRegressor(5, weights='uniform')

knn.fit(X, y)

y_test = knn.predict(X_test)

plot_prediction(X,y,X_test,y_test, "KNeighborsRegressor (k = 5, weights =_
↪ 'uniform')");
```

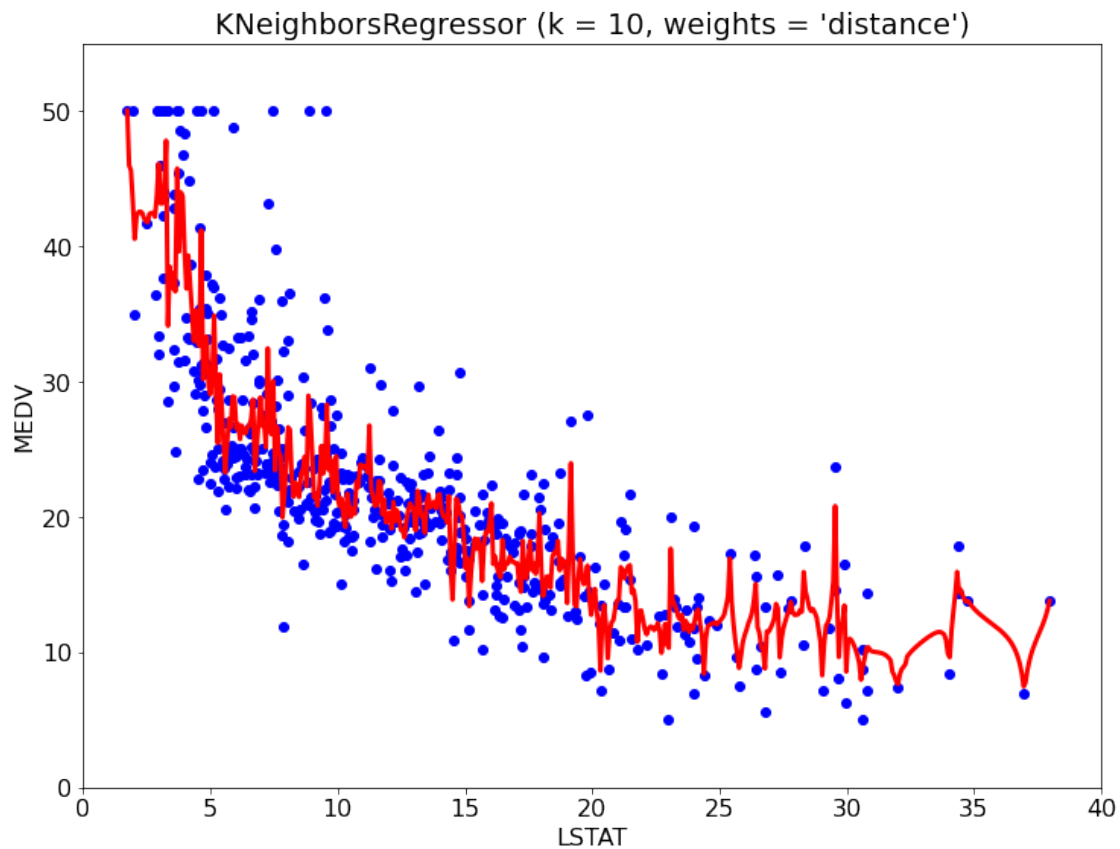


```
[ ]: knn = neighbors.KNeighborsRegressor(10, weights='distance')

knn.fit(X, y)


y_test = knn.predict(X_test)

plot_prediction(X,y,X_test,y_test, "KNeighborsRegressor (k = 10, weights =_
↪'distance')");
```



```
[ ]: regression_tree = DecisionTreeRegressor(max_depth=2)
      regression_tree.fit(X,y.ravel())
      y_test = regression_tree.predict(X_test)

      plot_prediction(X,y,X_test,y_test, "Regression Tree (Max Depth=2)");
```




```
[ ]: # implementation of the model developed using KNIME

def TreeModel(x):
    if (x<=9.725):
        if (x<=4.65):
            return 39.718
        else:
            return 26.6463
    else:
        if (x<=16.08):
            return 20.302
        else:
            return 14.2618

y_test = np.vectorize(TreeModel)(X_test)

plot_prediction(X,y,X_test,y_test, "KNIME Regression Tree");
```

```
[ ]: adaboost_model = AdaBoostRegressor(DecisionTreeRegressor(max_depth=2),
    ↪n_estimators=300, random_state=random_seed)
adaboost_model.fit(X, y.ravel())

y_test = adaboost_model.predict(X_test)

plot_prediction(X, y, X_test, y_test, "Adaboost (300 estimators Max Depth of
    ↪2)");
```

