

HousePricesPart2

October 14, 2021

1 House Prices: Advanced Regression Techniques (Part 2)

We now move to the actual model building part. As usual, we first import some of the libraries we will be needing.

```
[ ]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib

import matplotlib.pyplot as plt
from scipy import stats
from scipy.stats import skew
from scipy.stats import norm
from scipy.stats.stats import pearsonr

%config InlineBackend.figure_format = 'retina' #set 'png' here when working on_
↪notebook
%matplotlib inline
```

1.1 Read the Cleaned Data

First we load the data we cleaned in the previous part.

```
[ ]: X_train_complete = pd.read_csv("HousePricesTrainClean.csv")
X_test = pd.read_csv("HousePricesTestClean.csv")
X_train_complete.head()
```

```
[ ]: Unnamed: 0  MSSubClass  LotFrontage  LotArea  OverallQual  OverallCond  \
0            0      4.110874    4.189655  9.042040             7             5
1            1      3.044522    4.394449  9.169623             6             8
2            2      4.110874    4.234107  9.328212             7             5
3            3      4.262680    4.110874  9.164401             7             5
4            4      4.110874    4.442651  9.565284             8             5

      YearBuilt  YearRemodAdd  MasVnrArea  BsmtFinSF1  ...  SaleType_New  \
0          2003          2003    5.283204    6.561031  ...             0
1          1976          1976    0.000000    6.886532  ...             0
```

2	2001	2002	5.093750	6.188264	...	0
3	1915	1970	0.000000	5.379897	...	0
4	2000	2000	5.860786	6.486161	...	0

	SaleType_Oth	SaleType_WD	SaleCondition_Abnorml	SaleCondition_AdjLand	\
0	0	1	0	0	
1	0	1	0	0	
2	0	1	0	0	
3	0	1	1	0	
4	0	1	0	0	

	SaleCondition_Alloca	SaleCondition_Family	SaleCondition_Normal	\
0	0	0	1	
1	0	0	1	
2	0	0	1	
3	0	0	0	
4	0	0	1	

	SaleCondition_Partial	SalePrice
0	0	12.247699
1	0	12.109016
2	0	12.317171
3	0	11.849405
4	0	12.429220

[5 rows x 302 columns]

```
[ ]: X_train = X_train_complete.loc[:, 'MSSubClass': 'SaleCondition_Partial']
y = X_train_complete['SalePrice']
```

1.2 Linear Regression Models

Now we are going to use plain linear regression and regularized models (both L_1 Lasso and L_2 Ridge) from the scikit learn module. I'll also define a function that returns the cross-validation rmse error so we can evaluate our models and pick the best tuning par

```
[ ]: from sklearn.linear_model import LinearRegression, Ridge, RidgeCV, ElasticNet, \
      ↳Lasso, LassoCV
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error

def r2_cv(model, X_train, y, random_state=12345678):
    r2= cross_val_score(model, X_train, y, scoring="r2", cv =KFold(10, \
    ↳shuffle=True, random_state=random_state))
```

```
model_simple = LinearRegression()
model_simple.fit(X_train, y)
yp = model_simple.predict(X_train)

# compute R2 for train and using crossvalidation
r2_simple_train = r2_score(y, yp)
r2_xval_simple = r2_cv(model_simple, X_train, y)

# compute RMSE for train and using crossvalidation
rmse_simple_train = mean_squared_error(y, yp, multioutput='raw_values')
rmse_xval_simple = rmse_cv(model_simple, X_train, y)

print("Linear Regression")
print("====================")
print("\t\t\t\t\tTrain R2=%.3f"%(r2_simple_train))
print("\t10-fold Crossvalidation R2=%.3f"%(r2_xval_simple.mean()))
print("\t\t\t\t\tTrain RMSE=%.3f"%(rmse_simple_train))
print("\t10-fold Crossvalidation RMSE=%.3f"%(rmse_xval_simple.mean()))
```

1.3 Ridge Regression

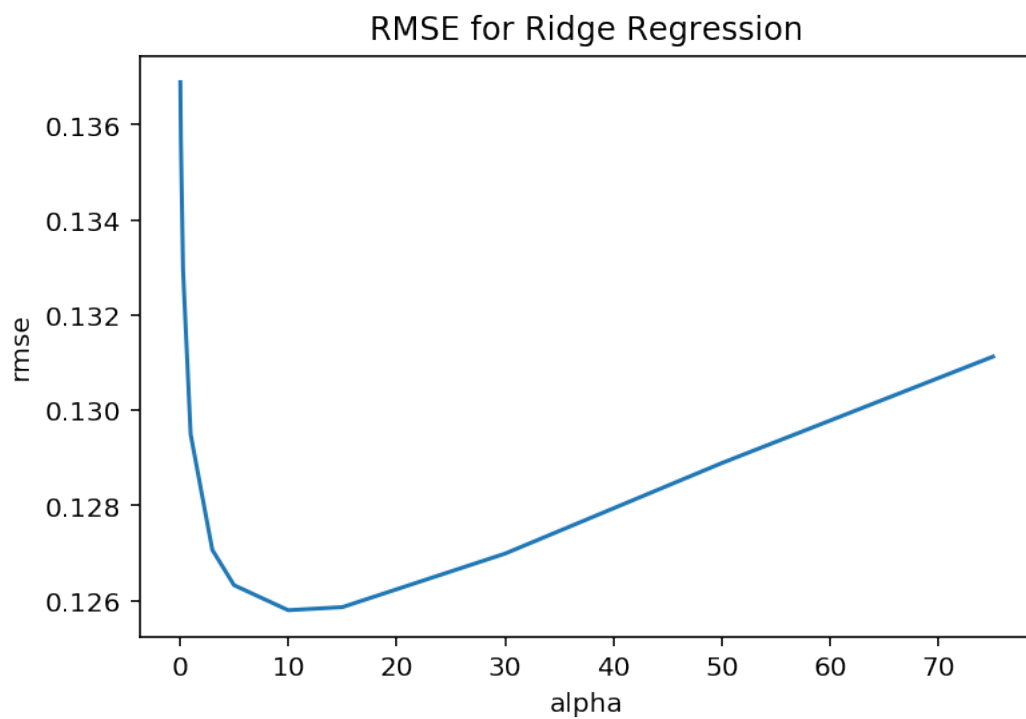
```
[ ]: alphas = [0.05, 0.1, 0.3, 1, 3, 5, 10, 15, 30, 50, 75]
cv_ridge = [rmse_cv(Ridge(alpha = alpha), X_train, y).mean() for alpha in
             ↪ alphas]
cv_r2_ridge = [r2_cv(Ridge(alpha = alpha), X_train, y).mean() for alpha in
               ↪ alphas]
```

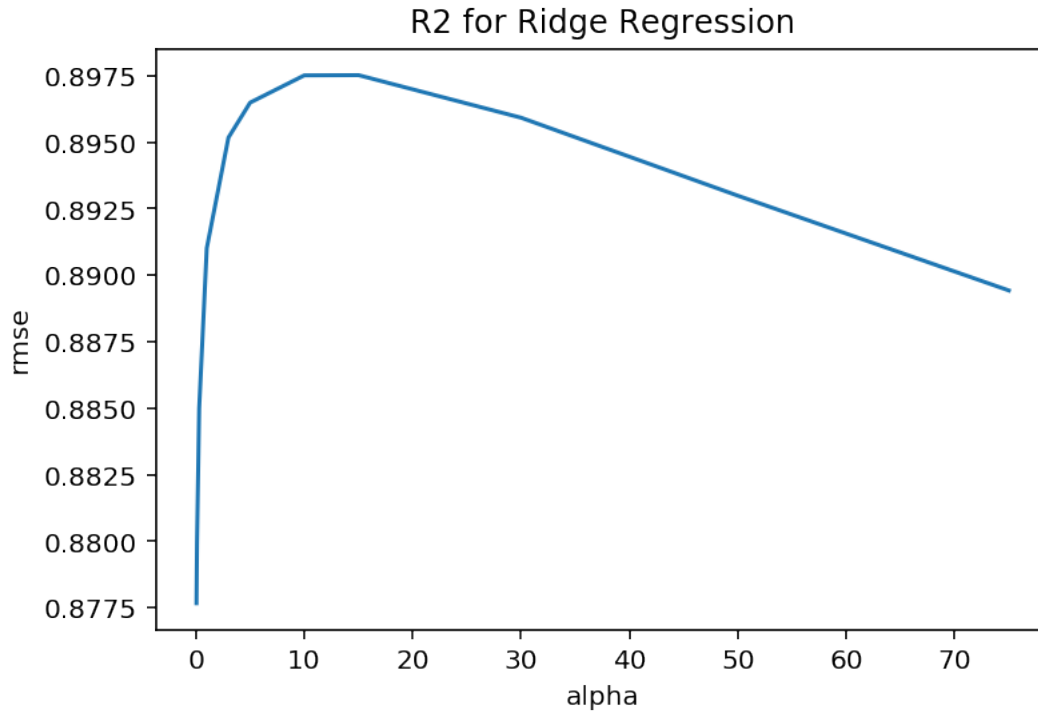
3

```
[ ]: cv_ridge = pd.Series(cv_ridge, index = alphas)
      cv_r2_ridge = pd.Series(cv_r2_ridge, index = alphas)

      cv_ridge.plot(title = "RMSE for Ridge Regression")
      plt.xlabel("alpha")
      plt.ylabel("rmse")
      plt.show()

      cv_r2_ridge.plot(title = "R2 for Ridge Regression")
      plt.xlabel("alpha")
      plt.ylabel("rmse")
      plt.show()
```





Note the U-ish shaped curve above. When alpha is too large the regularization is too strong and the model cannot capture all the complexities in the data. If however we let the model be too flexible (alpha small) the model begins to overfit. A value of $\alpha = 10$ is about right based on the plot above.

```
[ ]: print("Best RMSE %.3f for alpha %.3f"%(cv_ridge.min(),cv_ridge.idxmin()))
      print("Best R2 %.3f for alpha %.3f"%(cv_r2_ridge.max(),cv_r2_ridge.idxmax()))
      print("Why the difference? R2[10]=%.6f R2[15]=%.6f"%(cv_r2_ridge[cv_ridge.
        ↳idxmin()],cv_r2_ridge[cv_r2_ridge.idxmax()])))
```

Best RMSE 0.126 for alpha 10.000

Best R2 0.898 for alpha 15.000

Why the difference? R2[10]=0.897505 R2[15]=0.897513

So for the Ridge regression we get a rmse of about 0.126

1.4 Lasso Regression (L_1)

We now test Lasso regression. As before we need to test different values of α . Let's try out the Lasso model. For some reason the alphas in Lasso CV are really the inverse or the alphas in Ridge.

```
[ ]: alphas = [1, 0.1, 0.001, 0.0005]
      cv_lasso = [rmse_cv(Lasso(alpha = alpha), X_train, y).mean() for alpha in
        ↳alphas]
```

```

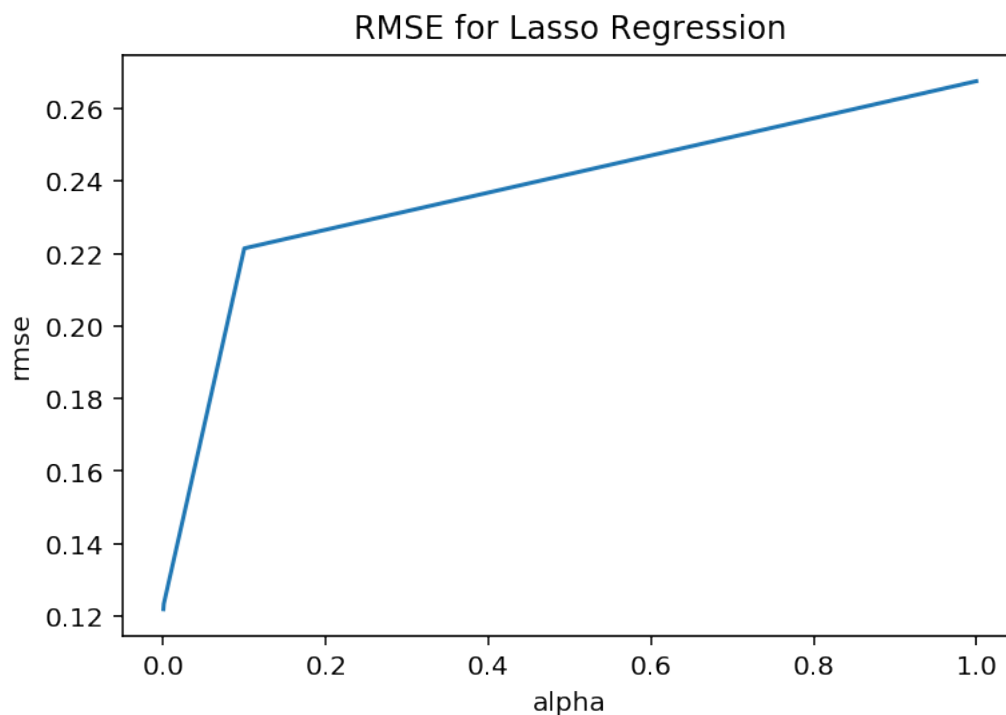
cv_r2_lasso = [r2_cv(Lasso(alpha = alpha), X_train, y).mean() for alpha in
    ↪ alphas]

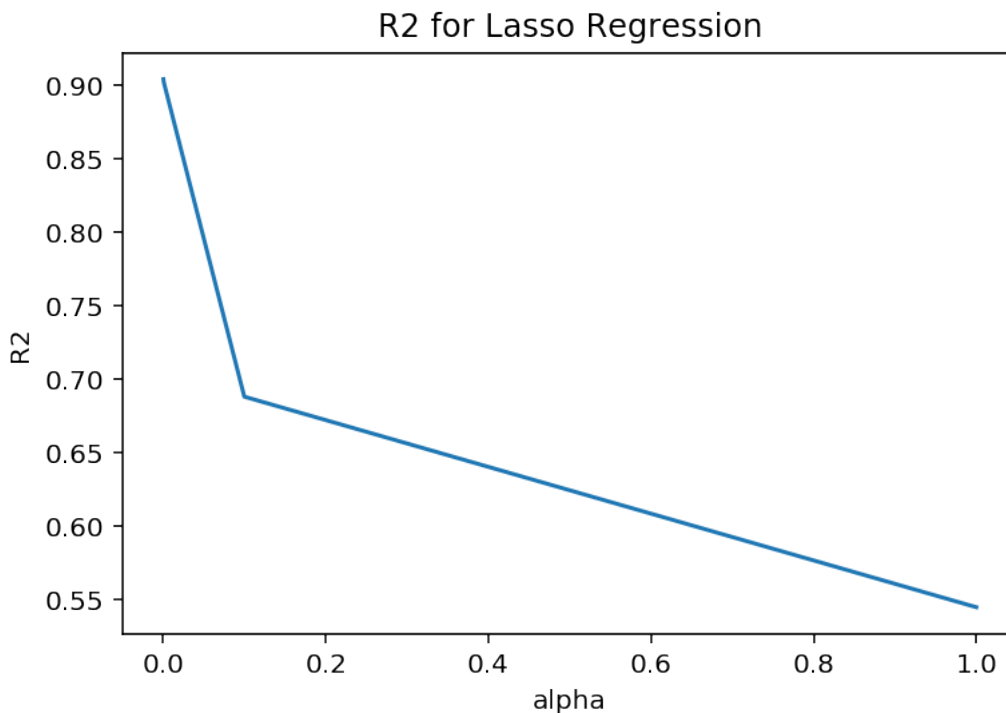
cv_lasso = pd.Series(cv_lasso, index = alphas)
cv_r2_lasso = pd.Series(cv_r2_lasso, index = alphas)

plt.figure(figsize=(6,4))
cv_lasso.plot(title = "RMSE for Lasso Regression")
plt.xlabel("alpha")
plt.ylabel("rmse")
plt.show()

plt.figure(figsize=(6,4))
cv_r2_lasso.plot(title = "R2 for Lasso Regression")
plt.xlabel("alpha")
plt.ylabel("R2")
plt.show()

```





```
[ ]: print("Best RMSE %.3f for alpha %.3f"%(cv_lasso.min(),cv_lasso.idxmin()))
      print("Best R2 %.3f for alpha %.3f"%(cv_r2_lasso.max(),cv_r2_lasso.idxmax()))
```

Best RMSE 0.122 for alpha 0.001

Best R2 0.904 for alpha 0.001

1.5 Lasso & Ridge with Built-in Crossvalidation

We performed the cross validation explicitly, but sklearn also provides two functions that include crossvalidation as part of the process, namely, LassoCV and RidgeCV.

```
[ ]: model_ridge = RidgeCV(alphas = [0.05, 0.1, 0.3, 1, 3, 5, 10, 15, 30, 50, 75],
      ↪cv=KFold(10, shuffle=True, random_state=12345678)).fit(X_train, y)
      model_lasso = LassoCV(alphas = [1, 0.1, 0.001, 0.0005],cv=KFold(10,
      ↪shuffle=True, random_state=12345678)).fit(X_train, y)
```

We evaluate the resulting model using the same procedure we applied before.

```
[ ]: rmse_lasso2 = rmse_cv(model_lasso, X_train, y)
      rmse_ridge2 = rmse_cv(model_ridge, X_train, y)

      r2_lasso2 = r2_cv(model_lasso, X_train, y)
      r2_ridge2 = r2_cv(model_ridge, X_train, y)
```

Finally, we compare the results achieved with this procedure with the ones compared with the

explicit cross validation.

```
[ ]: print("Ridge Regression (10-fold crossvalidation)")
print("\tRMSE=%.3f R2=%.3f for Alpha=%.3f"%(rmse_ridge2.mean(), r2_ridge2.
    ↪mean(), model_ridge.alpha_))
print("\n")
print("Lasso Regression (10-fold crossvalidation)")
print("\tRMSE=%.3f R2=%.3f for Alpha=%.3f"%(rmse_lasso2.mean(), r2_lasso2.
    ↪mean(), model_lasso.alpha_))
```

```
Ridge Regression (10-fold crossvalidation)
      RMSE=0.126 R2=0.896 for Alpha=15.000
```

```
Lasso Regression (10-fold crossvalidation)
      RMSE=0.122 R2=0.904 for Alpha=0.001
```

The results are similar to the ones we obtained with the explicit search for α .

1.6 Feature Selection

One nice feature of Lasso regression is that by zeroing out the coefficients of variables that it deems unimportant, it actually performs feature selection. In practise, it simplifies the data while building the model.

```
[ ]: coef = pd.Series(model_lasso.coef_, index = X_train.columns)

[ ]: print("Lasso picked " + str(sum(coef != 0)) + " variables and eliminated the_
    ↪other " + str(sum(coef == 0)) + " variables")
```


Lasso picked 105 variables and eliminated the other 195 variables

Note that, the process is stochastic and thus we cannot be sure that the selected variables are exactly the good ones so one approach is to run the procedure several times (using bootstrap which we will discuss later in the course) to check how robust the feature selection is. For instance, we check if there are variables whose weight is always zeroed out. As an example, let's plot the largest and the smallest coefficients.

```
[ ]: imp_coef = pd.concat([coef.sort_values().head(20),
                          coef.sort_values().tail(20)])

[ ]: matplotlib.rcParams['figure.figsize'] = (8.0, 10.0)
imp_coef.plot(kind = "barh")
plt.title("Coefficients in the Lasso Model")
plt.xlim(-0.7,0.5)
```

```
[ ]: (-0.7, 0.5)
```


GrLivArea appears to be the most important positive feature (it identifies the above ground area by area square feet so it makes sense). Then a few other location and quality features contributed positively. Some of the negative features make less sense and would be worth looking into more.

Let's also check the residuals, that is the difference between the predicted and the actual value plotted against the predicted value.

```
[ ]: matplotlib.rcParams['figure.figsize'] = (6.0, 6.0)

preds = pd.DataFrame({"preds":model_lasso.predict(X_train), "true":y})
preds["residuals"] = preds["true"] - preds["preds"]
preds.plot(x = "preds", y = "residuals",kind = "scatter")
plt.xlabel("Predicted Value")
plt.ylabel("Residual Value")
```

```
[ ]: Text(0,0.5,'Residual Value')
```



```
[ ]: ridge_coef = pd.Series(model_ridge.coef_, index = X_train.columns)
print("Ridge picked " + str(sum(ridge_coef != 0)) + " variables and eliminated_
↳the other " + str(sum(ridge_coef == 0)) + " variables")
```

Ridge picked 300 variables and eliminated the other 0 variables

```
[ ]: ridge_imp_coef = pd.concat([ridge_coef.sort_values().head(20),
                                ridge_coef.sort_values().tail(20)])
matplotlib.rcParams['figure.figsize'] = (8.0, 10.0)
ridge_imp_coef.plot(kind = "barh")
plt.title("Coefficients in the Ridge Model")
plt.xlim(-0.7,0.5)
```

```
[ ]: (-0.7, 0.5)
```

